Exemplar Auditing Lifecycle

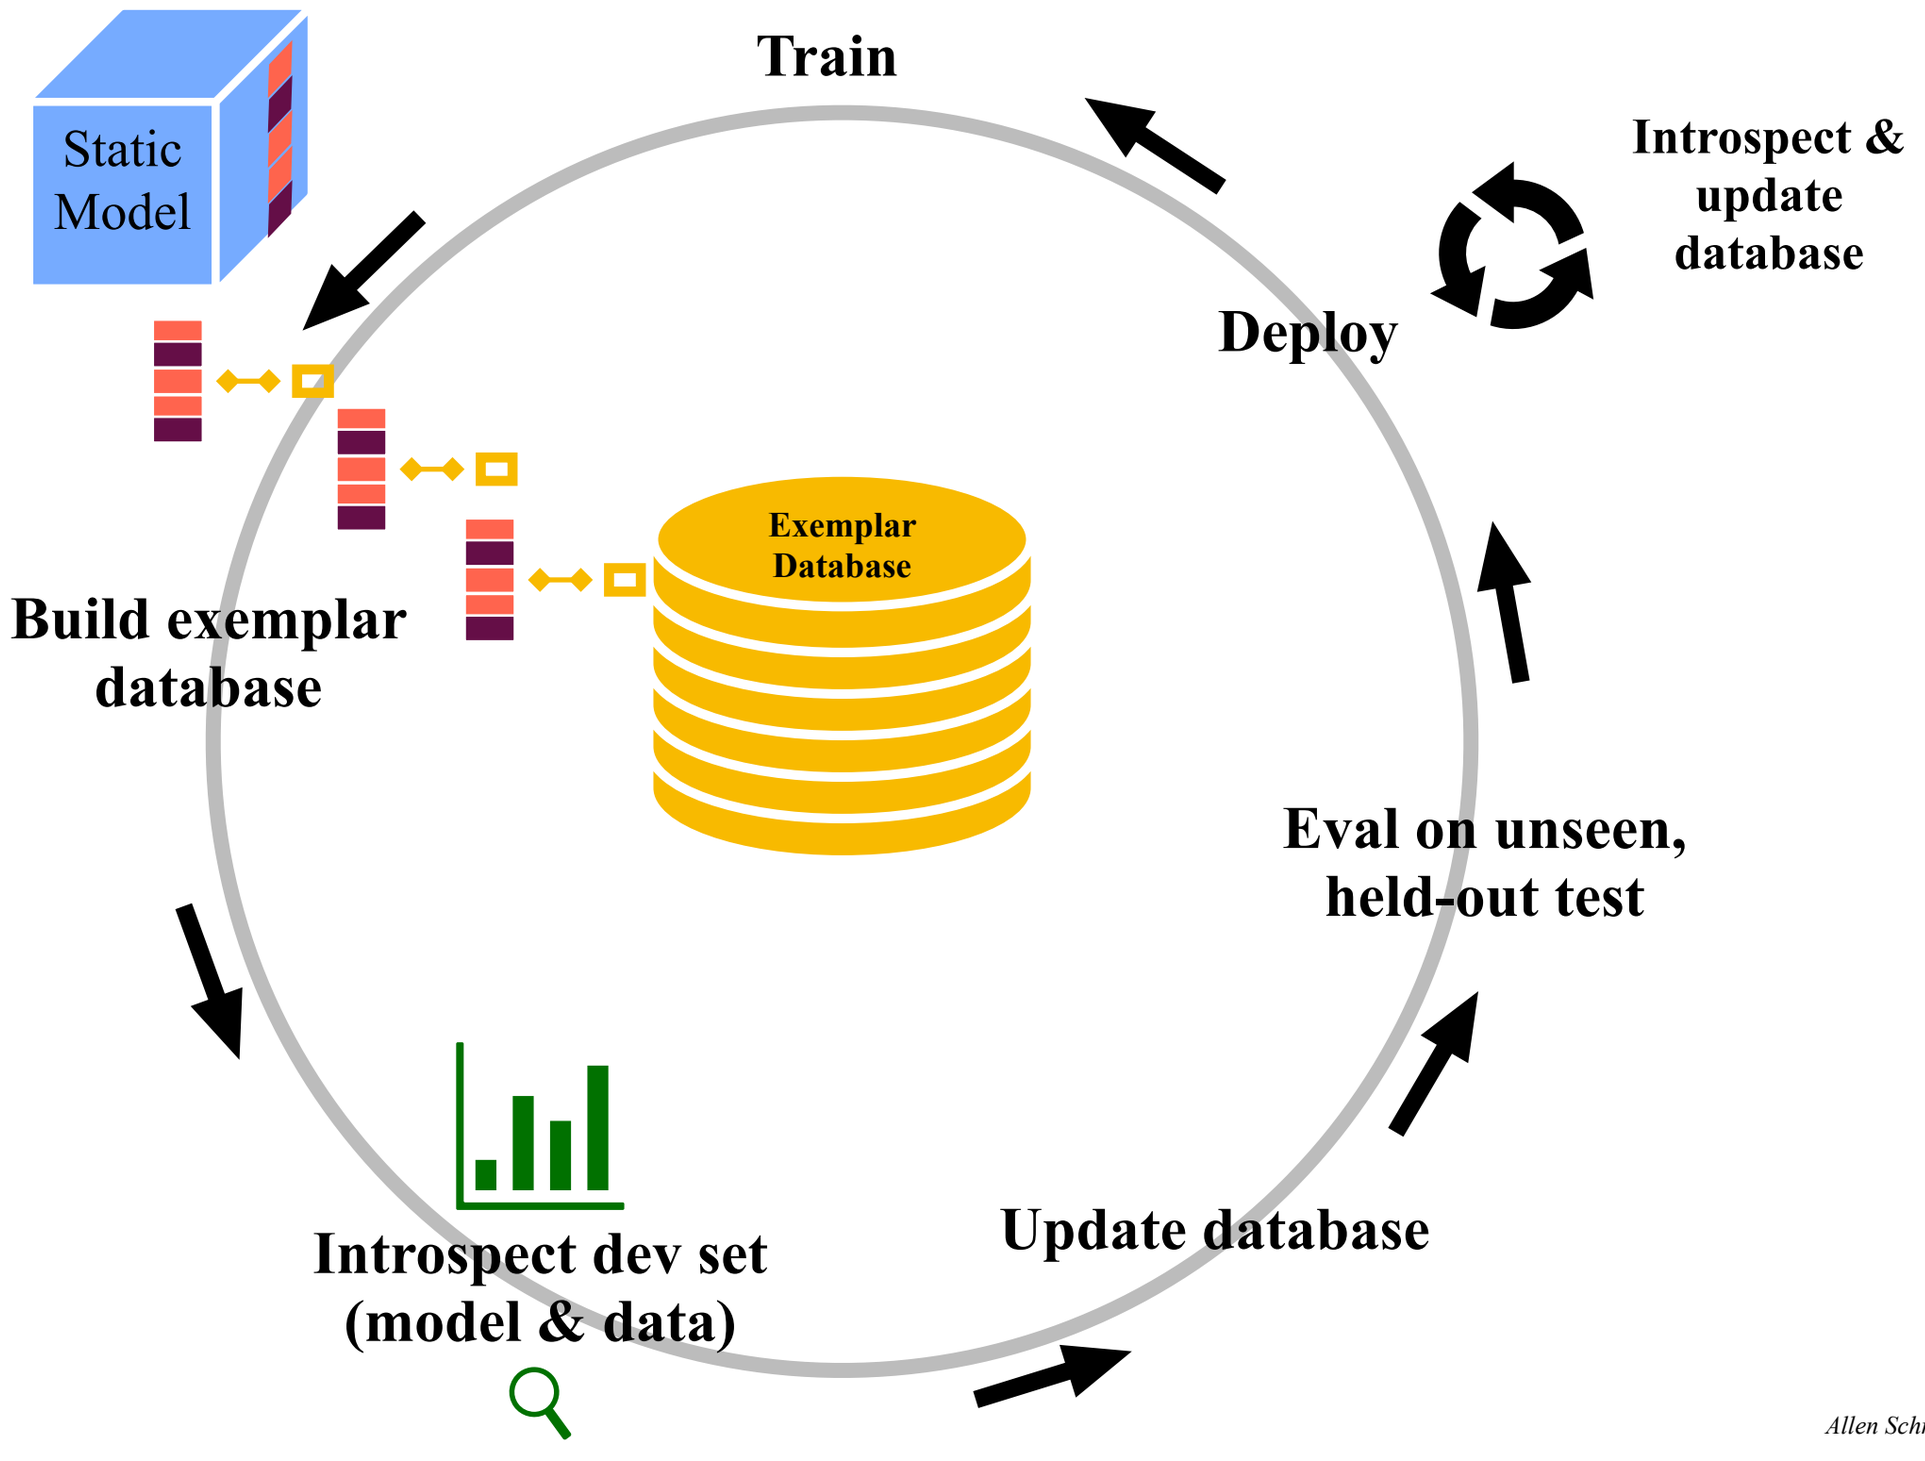

Allen Schmaltz, 2019-2021

Out-of-Domain Settings

- Pre-train with as much data as possible
- Add as much data as possible to the database
    - Corral the in-domain space, around the ball of the observed data
    - Never predict over out-of-domain data in high-risk settings. Instead: Rearrange the deployment to handle non-admitted predictions.

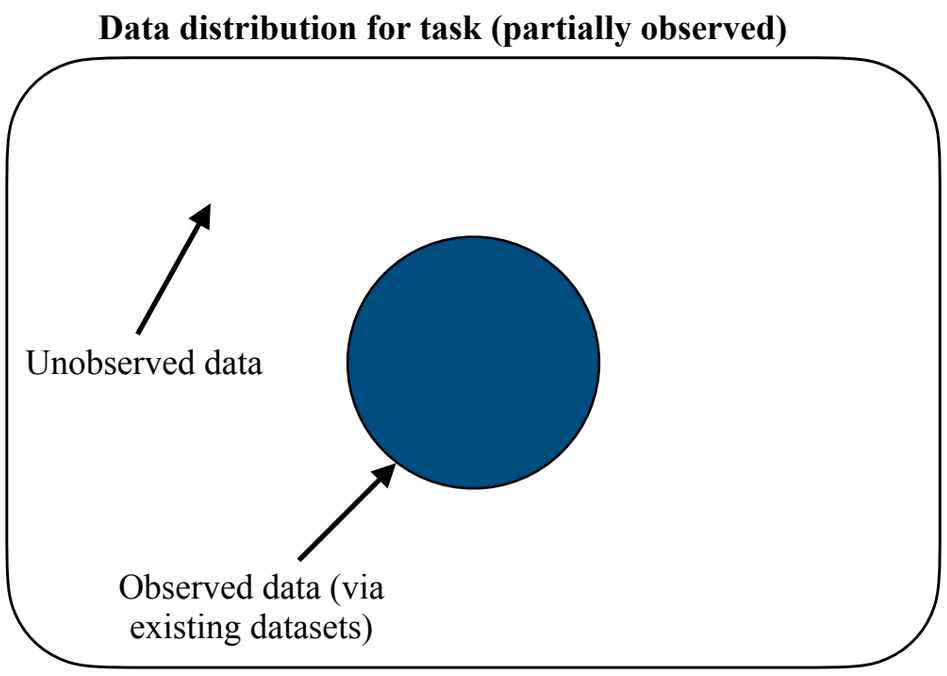

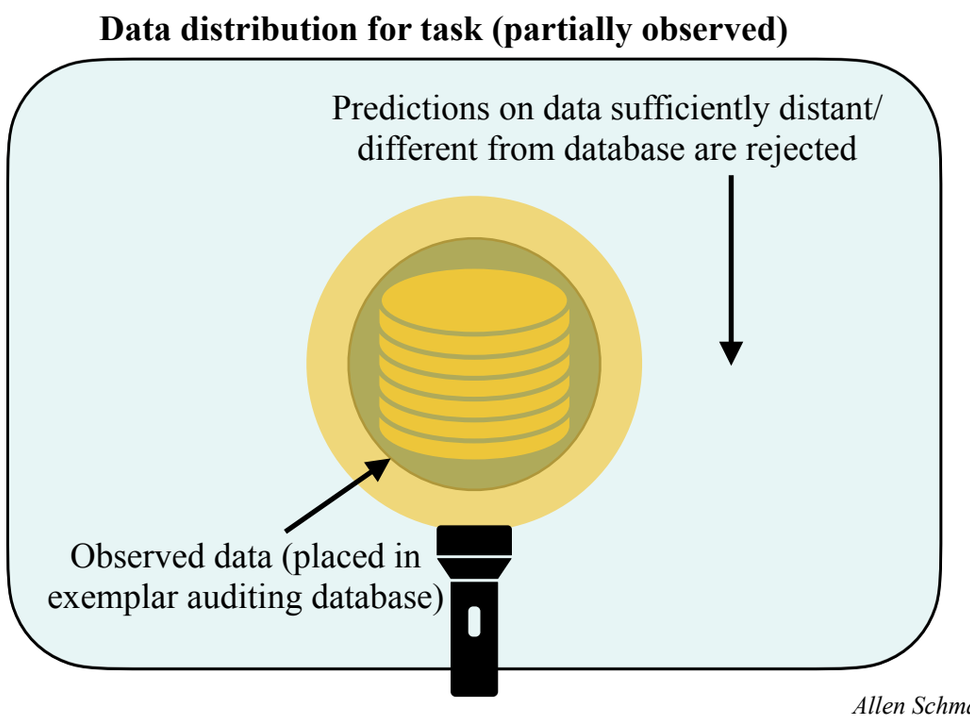

Allen Schmaltz, 2019-2021

Implementations

- Binary classification: $f : X \rightarrow \{0,1\}$

  Unique side effect: Sequence labeling: $f : X \rightarrow Y_1, \ldots, Y_{|x|}$

  - "Detecting Local Insights from Global Labels: Supervised & Zero-Shot Sequence Labeling via a Convolutional Decomposition"

- Multi-label classification: $f : X \rightarrow 2^{|Y|}$

  - "Exemplar Auditing for Multi-Label Biomedical Text Classification"

- Retrieval-classification: $f : X \times \mathscr{D} \rightarrow \left\langle \{0,1,2\},\ 2^{|D|} \right\rangle$

  - "Coarse-to-Fine Memory Matching for Joint Retrieval and Classification"

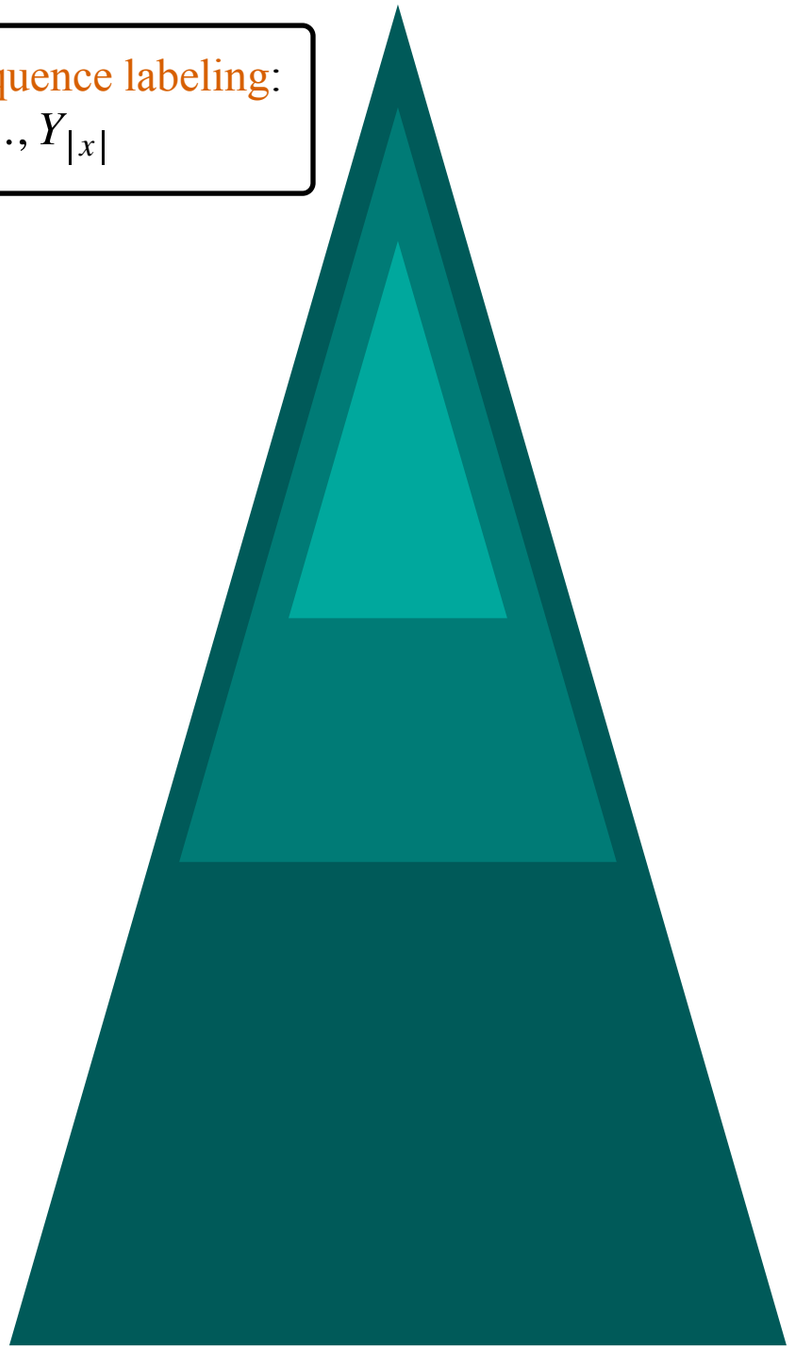

An End-to-End Retrieval-Classification Model via a Coarse-to-Fine Search over Dense Representations

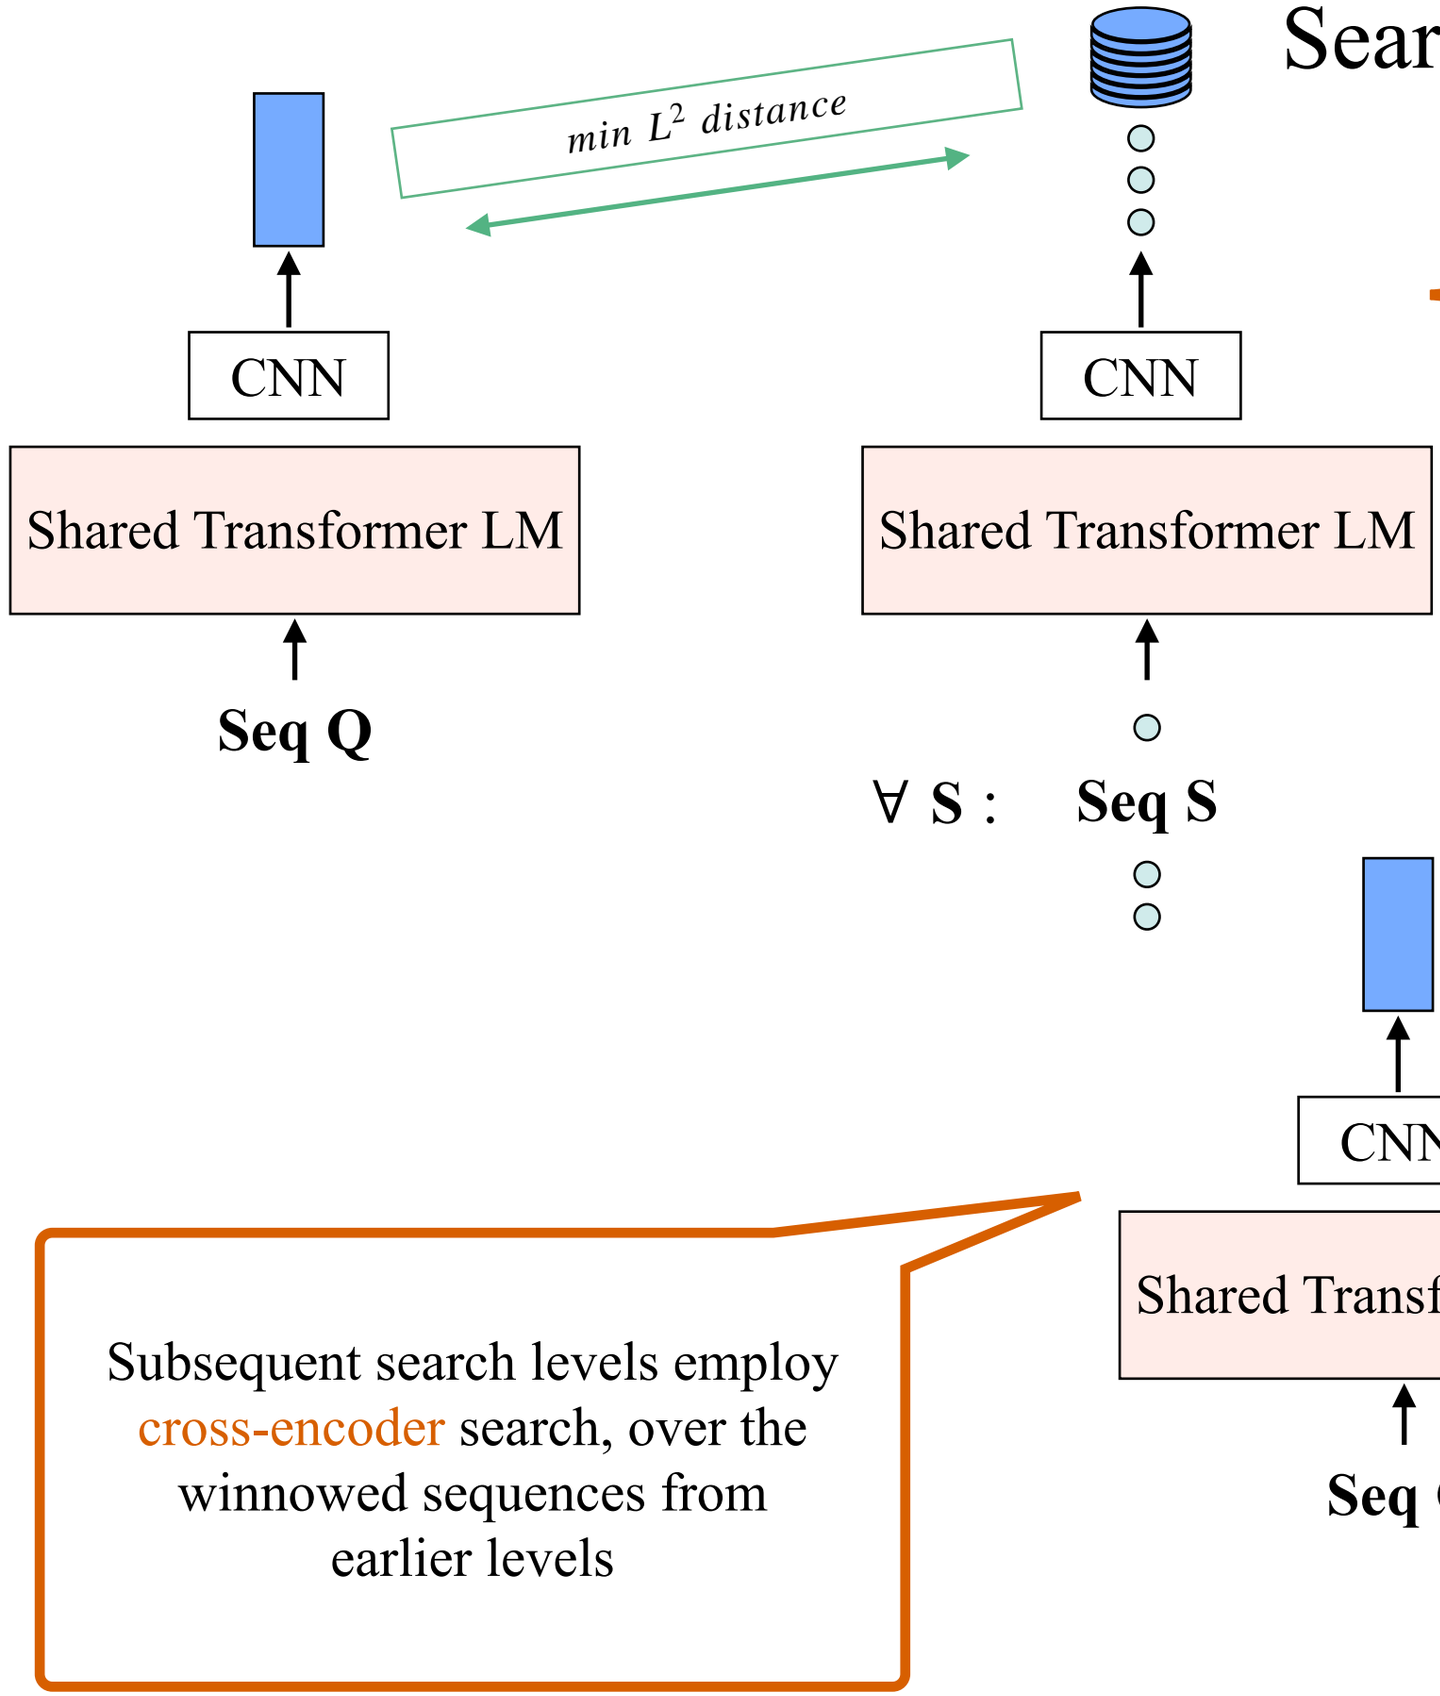

Initial, coarse bi-encoder search

Subsequent search levels employ cross-encoder search, over the winnowed sequences from earlier levels

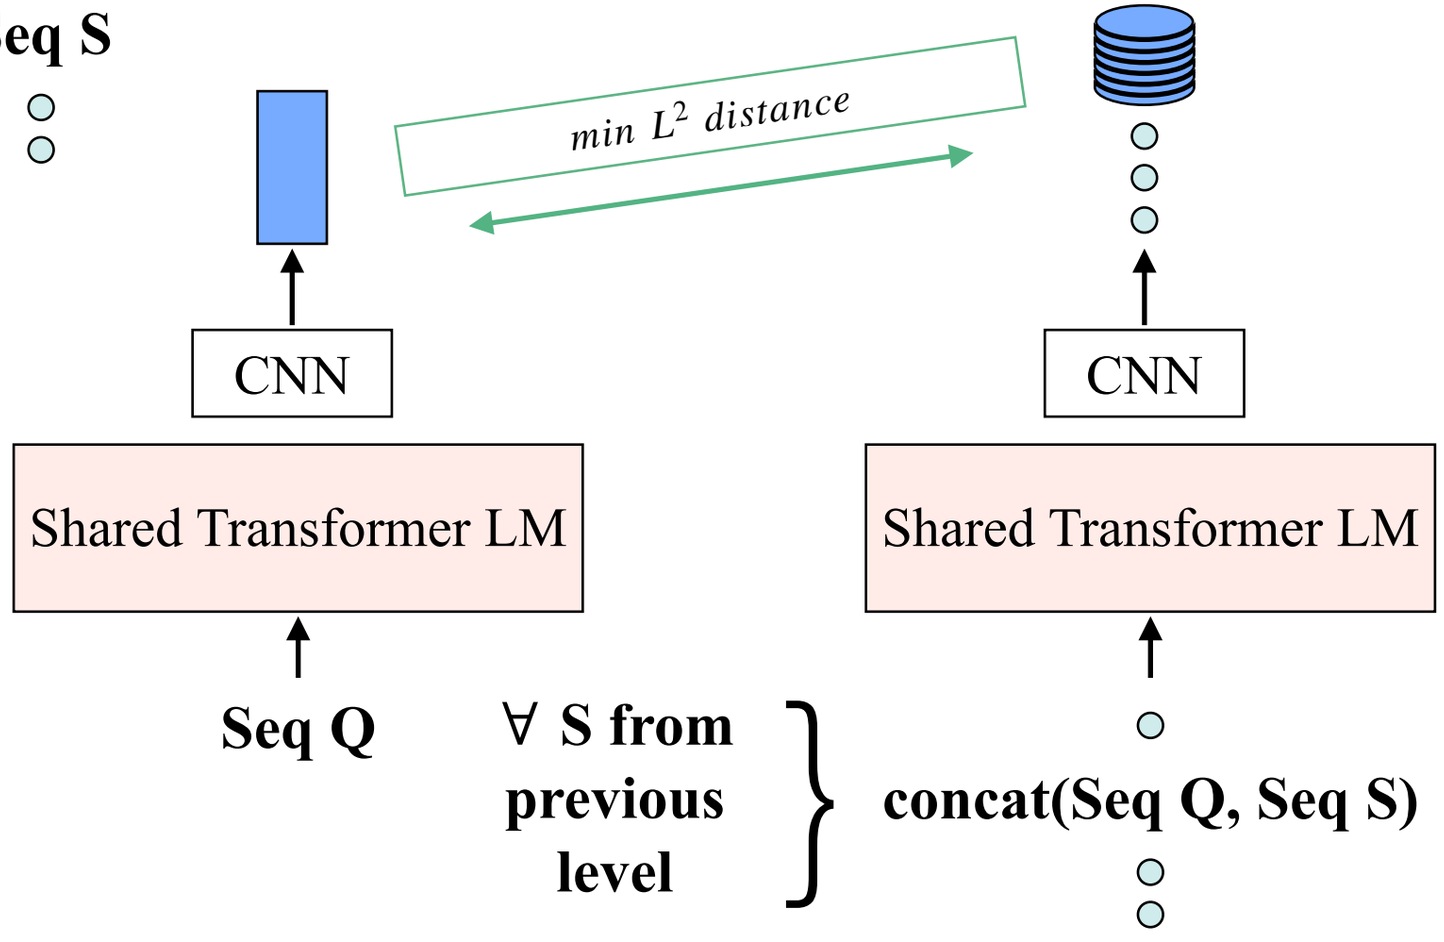

Joint Retrieval and Classification Training

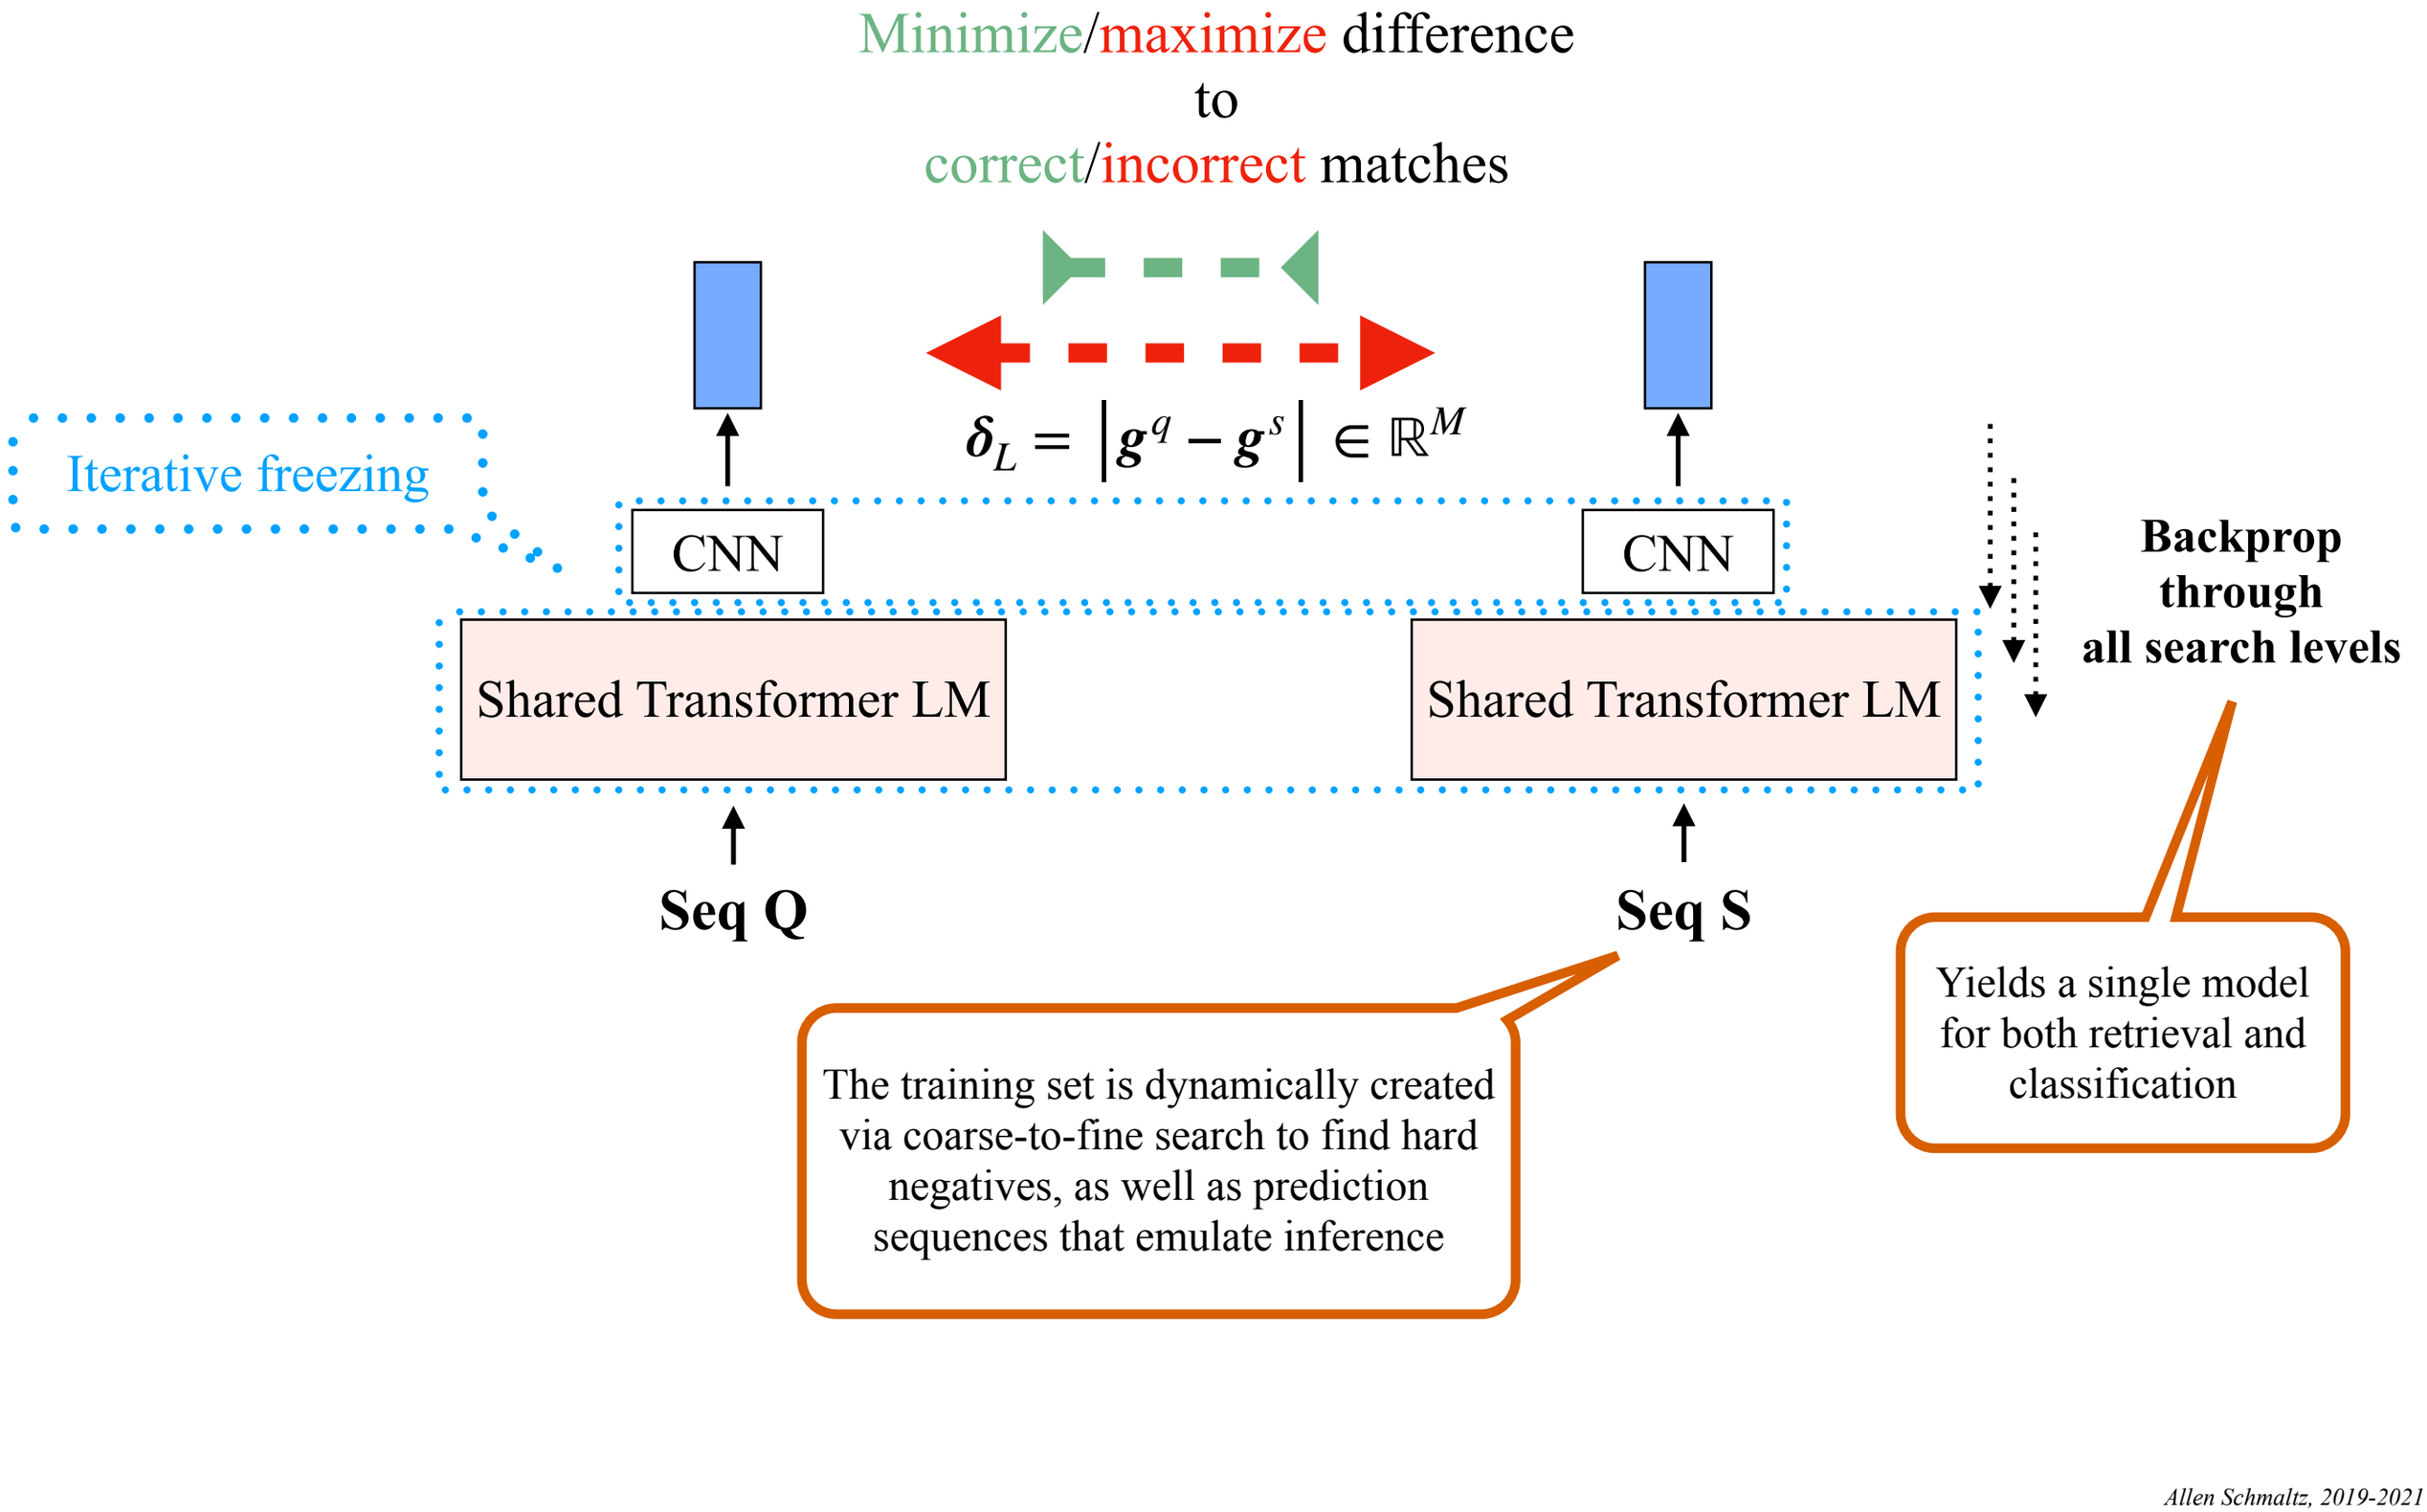

Multi-Sequence Representation Composition for Exemplar Auditing

Dense representation of query sequence:

$$\boldsymbol{g}^q \in \mathbb{R}^{1000} = \begin{bmatrix} g_1^q \\ g_2^q \\ \vdots \\ g_{1000}^q \end{bmatrix}$$

CNN

Shared Transformer LM

Seq Q

$$\boldsymbol{\delta}_{L_2} = \left| \boldsymbol{g}^q - \boldsymbol{g}^s \right| \in \mathbb{R}^M$$

CNN

Shared Transformer LM

concat(Seq Q, Seq S)

Dense representation of support sequence:

$$\boldsymbol{g}^s \in \mathbb{R}^{1000} = \begin{bmatrix} g_1^s \\ g_2^s \\ \vdots \\ g_{1000}^s \end{bmatrix}$$

Search levels

$$\boldsymbol{\delta}_{L_3} = \left| \boldsymbol{g}^q - \boldsymbol{g}^s \right| \in \mathbb{R}^M$$

CNN

Shared Transformer LM

Seq Q

CNN

Shared Transformer LM

concat(Label, Seq Q, Seq S*)

Final composed representation: concat($\boldsymbol{\delta}_{L_2}$, $\boldsymbol{\delta}_{L_3}$)

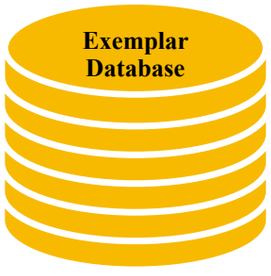

Token-Level Representations for Exemplar Auditing

Identify the dense representation of a token-level feature using [multi-] binary labeling via a convolutional decomposition (optionally, with priors to encourage/discourage particular features)

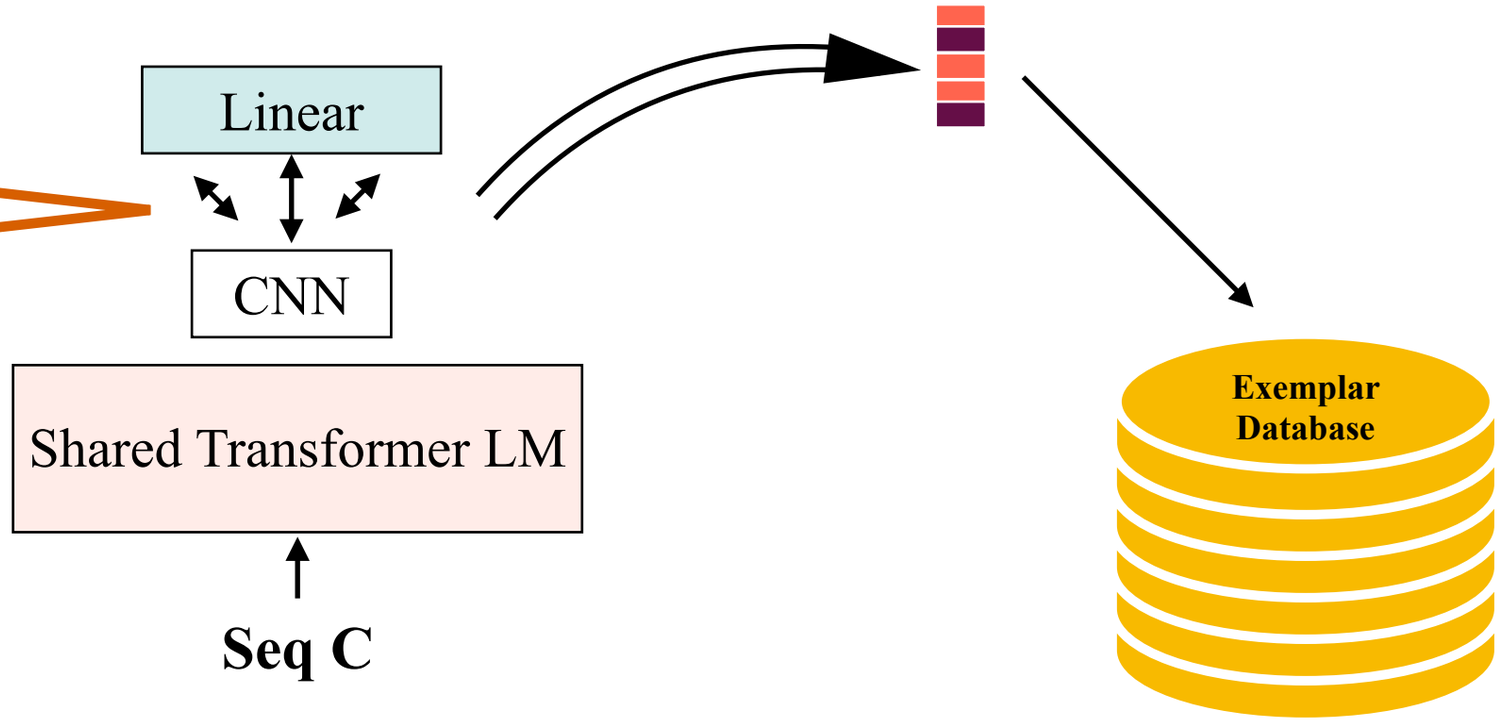

Prospective Outlook: Interlocking distance constraints across input modalities and tasks via a single, shared model and a dense database…

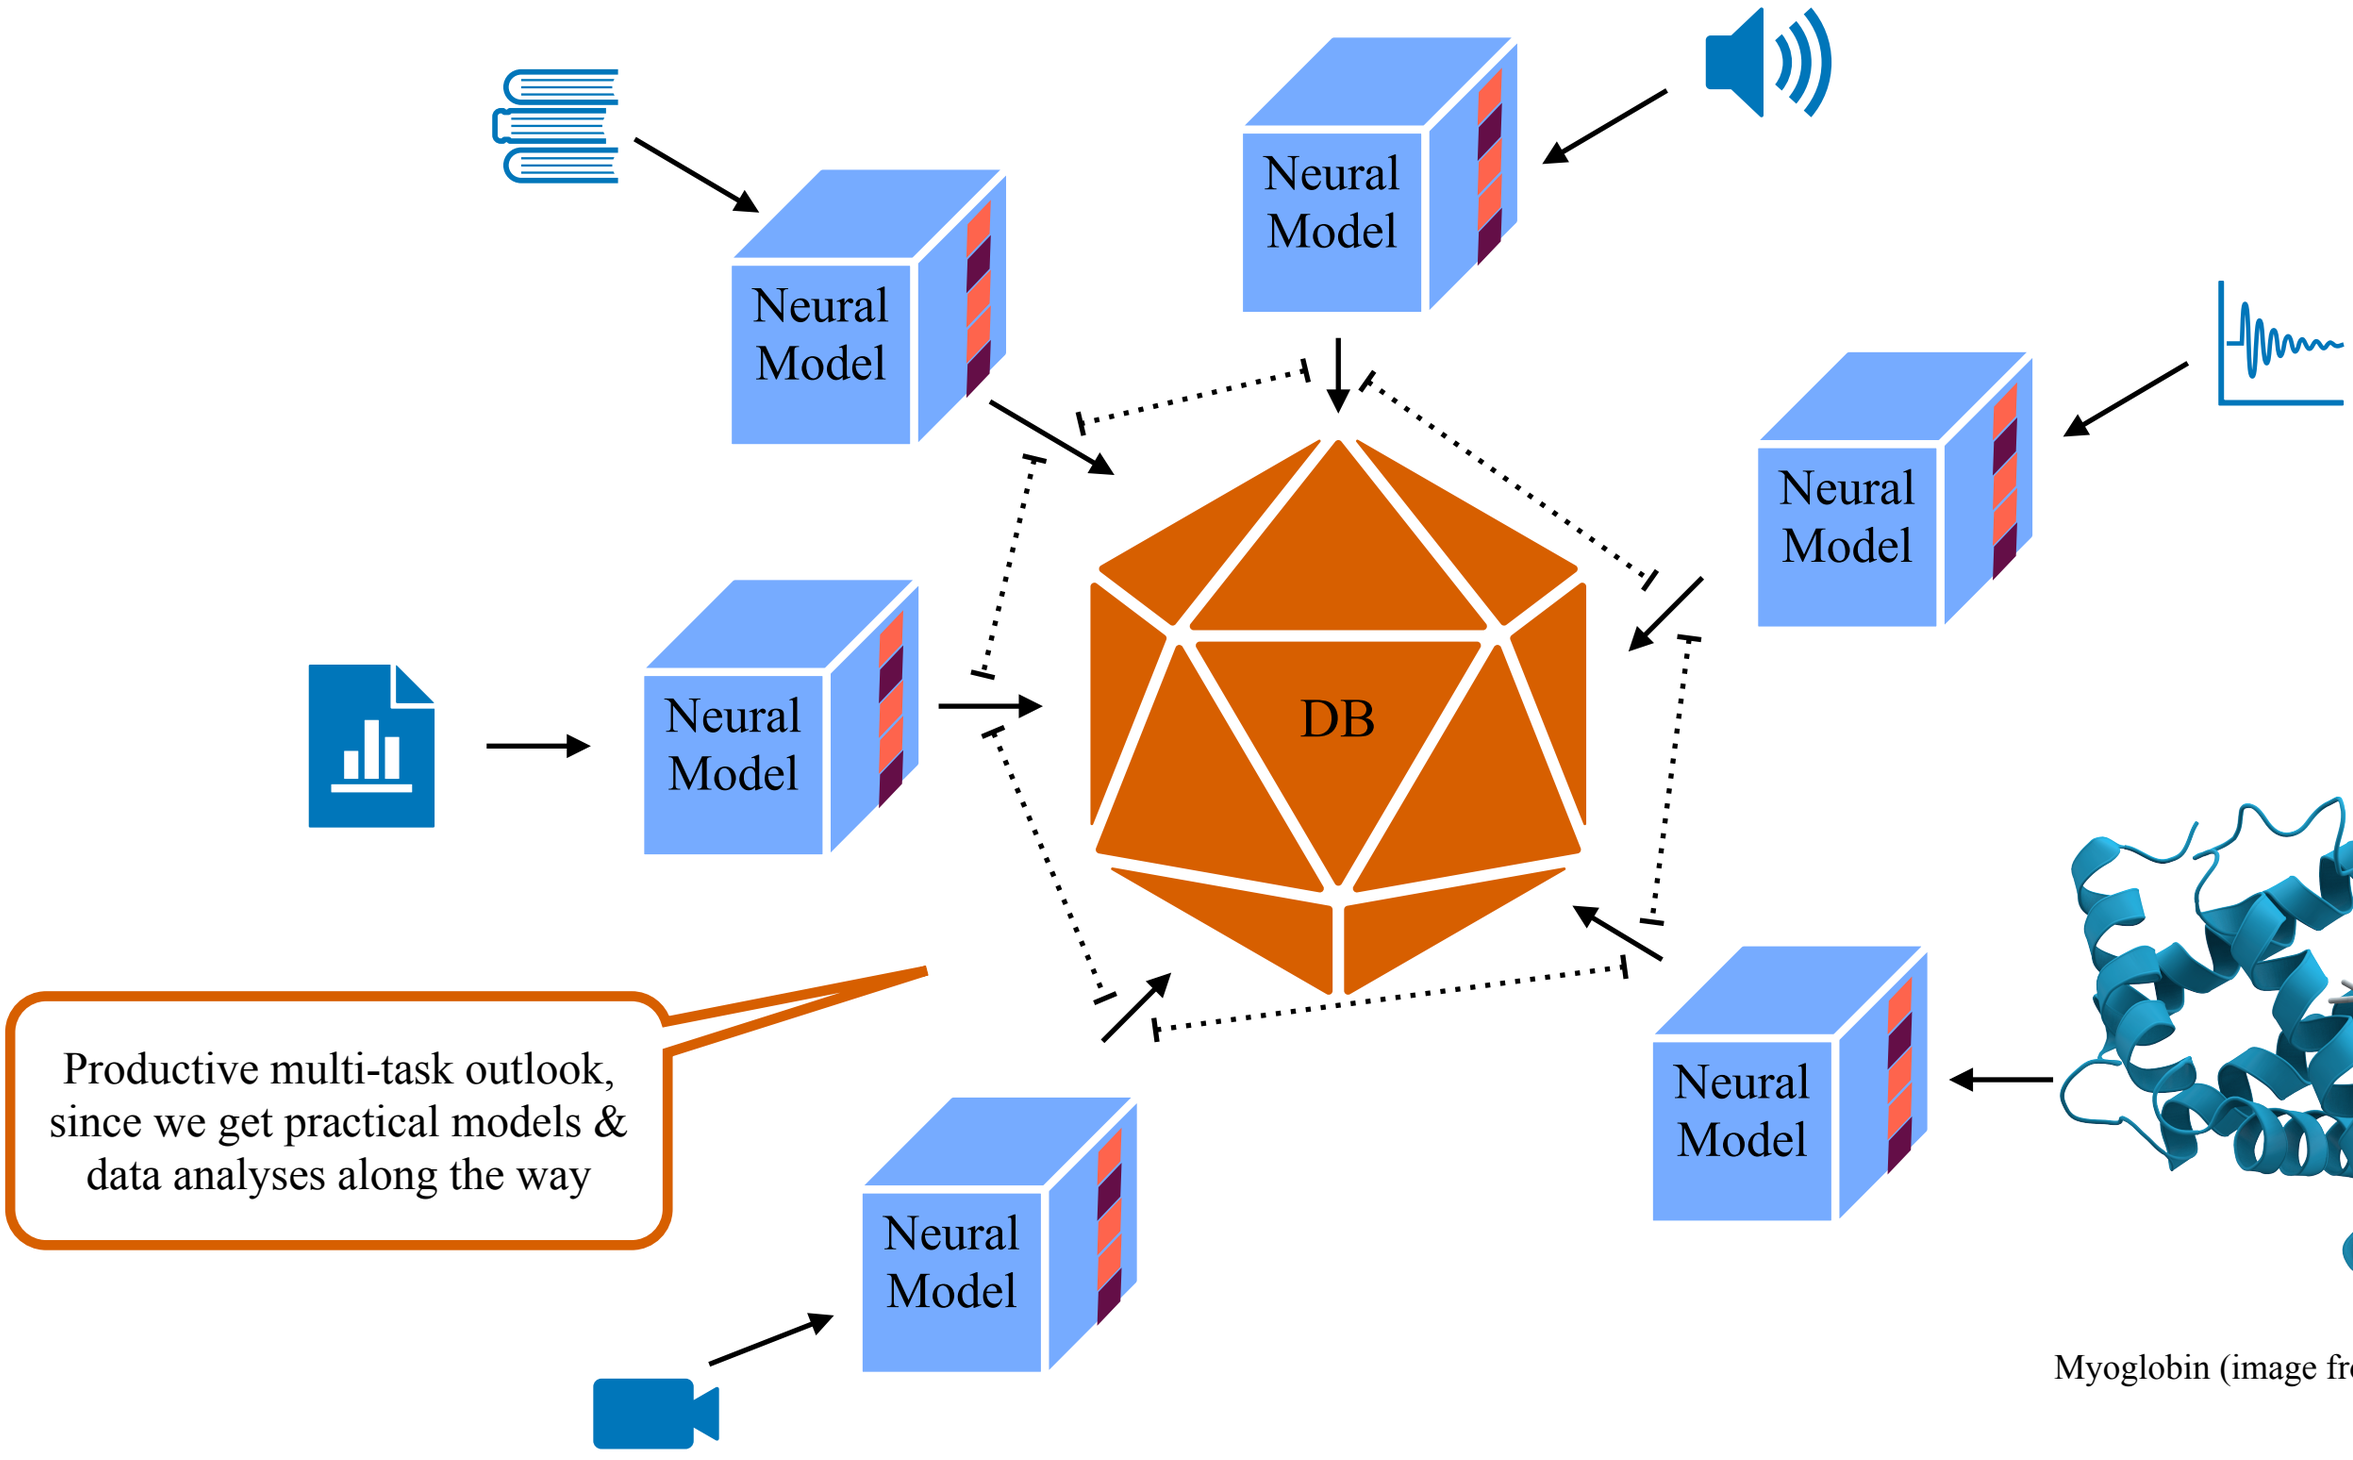

Myoglobin (image from Wikipedia)